

Causal Inference: Introduction to Propensity Score Matching

Joseph Kush

5/17/2022

Purpose:

This R script introduces propensity score matching methods for causal inference. Example school-level data is then generated, in which there is imbalance on baseline covariates between schools implementing treatment versus control. After examining the extent of imbalance, propensity scores are estimated and schools are matched. Finally, the effect of treatment on the outcome of interest (math) is estimated and compared to the known population ATT.

Data generation

First, load necessary libraries and set seed for reproducibility.

```
library("cobalt")
```

```
## Warning: package 'cobalt' was built under R version 4.0.2
```

```
## cobalt (Version 4.3.1, Build Date: 2021-03-30 09:50:18 UTC)
```

```
library("MatchIt")
```

```
## Warning: package 'MatchIt' was built under R version 4.0.2
```

```
##
```

```
## Attaching package: 'MatchIt'
```

```
## The following object is masked from 'package:cobalt':
```

```
##
```

```
## lalonde
```

```
library("survey")
```

```
## Warning: package 'survey' was built under R version 4.0.2
```

```
## Loading required package: grid
```

```
## Loading required package: Matrix
```

```
## Loading required package: survival
```

```
##
## Attaching package: 'survey'

## The following object is masked from 'package:graphics':
##
##      dotchart

set.seed(1234)
```

In this example, we generate school-level covariates such as enrollment, student:teacher ratio, and % of students receiving free and reduced-price meals for $N = 1,000$ schools.

```
# Generate school-level data
sample_size <- 1000
schID <- 1:sample_size

enrollment <- as.integer(runif(n = sample_size, min = 100, max = 1000))
st_ratio <- rnorm(n = sample_size, mean = 0.2, sd = 0.05)
susp <- runif(n = sample_size, min = 0, max = 0.2)
farms <- rnorm(n = sample_size, mean = st_ratio, sd = 0.03)
sped <- rnorm(n = sample_size, mean = 0.15, sd = 0.03)
minority <- runif(n = sample_size, min = 0, max = 0.8)
ell <- rnorm(n = sample_size, mean = farms, sd = 0.02)
disable <- rnorm(n = sample_size, mean = 0.15, sd = 0.03)
read <- rnorm(n = sample_size, mean = 0.8, sd = 0.05)
```

Next, a binary treatment indicator is generated for each school according to a treatment-selection model. Logit odds (x) are created with a binary treatment indicator generated from a binomial distribution with a probability of exposure equal to $\frac{e^x}{1+e^x}$.

```
# Create a binary treatment indicator
# First, create logit odds (hint: intercept sets treatment prevalence)
logit_treat <- 0.7 + 0*enrollment + 2.7*st_ratio + 1.9*susp + 1.3*farms +
  9*sped*minority + 1.8*ell + 2.6*disable + -3.7*read

# Next, convert logit odds into probability
prob_treat <- exp(logit_treat)/(1 + exp(logit_treat))

# Finally, generate binary treatment indicator from binomial
# distribution of 1 trial with P = prob_treat
treat <- rbinom(sample_size, 1, prob_treat)

# Create potential outcomes for reading scores
treatment_effect <- 2

math_0 <- (80 + 0*treatment_effect +
  0*enrollment + 2.7*st_ratio + 1.9*susp + 1.3*farms +
  9*sped*minority + 1.8*ell + 2.6*disable + -3.7*read +
  rnorm(n = sample_size, mean = 0, sd = 3)) / 100
```

```

math_1 <- (80 + 1*treatment_effect +
  0*enrollment + 2.7*st_ratio + 1.9*susp + 1.3*farms +
  9*sped*minority + 1.8*ell + 2.6*disable + -3.7*read +
  rnorm(n = sample_size, mean = 0, sd = 3)) / 100

math <- ifelse(treat == 1, math_1, math_0)

```

Everything is placed into a data.frame.

```

# Combine all variables into a data frame
df <- data.frame(cbind(schID, enrollment, st_ratio, susp, farms, sped,
  minority, ell, disable, read, treat, math))

summary(df)

```

```

##      schID      enrollment      st_ratio      susp
## Min.   : 1.0    Min.   :100.0    Min.   :0.05348    Min.   :0.0001799
## 1st Qu.:250.8    1st Qu.:332.0    1st Qu.:0.16839    1st Qu.:0.0500234
## Median :500.5    Median :558.5    Median :0.20025    Median :0.0955312
## Mean   :500.5    Mean   :556.1    Mean   :0.19942    Mean   :0.0991982
## 3rd Qu.:750.2    3rd Qu.:782.2    3rd Qu.:0.23108    3rd Qu.:0.1480819
## Max.   :1000.0    Max.   :999.0    Max.   :0.33529    Max.   :0.1998465
##      farms      sped      minority      ell
## Min.   :0.02307    Min.   :0.05871    Min.   :0.0008865    Min.   :0.01205
## 1st Qu.:0.16034    1st Qu.:0.12925    1st Qu.:0.2011978    1st Qu.:0.16036
## Median :0.20098    Median :0.15013    Median :0.3978692    Median :0.19845
## Mean   :0.20028    Mean   :0.14971    Mean   :0.4009244    Mean   :0.19980
## 3rd Qu.:0.24007    3rd Qu.:0.16981    3rd Qu.:0.5997060    3rd Qu.:0.24207
## Max.   :0.36934    Max.   :0.24360    Max.   :0.7969689    Max.   :0.37539
##      disable      read      treat      math
## Min.   :0.06121    Min.   :0.6400    Min.   :0.000    Min.   :0.7018
## 1st Qu.:0.13091    1st Qu.:0.7672    1st Qu.:0.000    1st Qu.:0.7826
## Median :0.14946    Median :0.8021    Median :1.000    Median :0.8036
## Mean   :0.15003    Mean   :0.8022    Mean   :0.513    Mean   :0.8036
## 3rd Qu.:0.16990    3rd Qu.:0.8365    3rd Qu.:1.000    3rd Qu.:0.8252
## Max.   :0.24885    Max.   :0.9809    Max.   :1.000    Max.   :0.8998

```

```
head(df)
```

```

##  schID enrollment st_ratio      susp      farms      sped minority      ell
## 1      1        202 0.2492390 0.19271987 0.2200244 0.1384478 0.1342445 0.1924882
## 2      2        660 0.1387631 0.04135248 0.1357742 0.1654217 0.3716408 0.1177858
## 3      3        648 0.2354863 0.01723949 0.2321643 0.1592401 0.1663240 0.2630359
## 4      4        661 0.1945390 0.04320560 0.2303048 0.2051746 0.6150317 0.1939220
## 5      5        874 0.2891304 0.04792931 0.2394538 0.1978021 0.3383099 0.2227598
## 6      6        676 0.1878278 0.03943219 0.1564585 0.1392541 0.6416494 0.1821018
##      disable      read treat      math
## 1 0.1362781 0.8050001      0 0.7831675
## 2 0.1479290 0.7342882      0 0.7797560
## 3 0.1410184 0.8387209      0 0.7617098
## 4 0.1362928 0.7748310      1 0.7868684
## 5 0.1708369 0.8553123      1 0.7726073
## 6 0.1542474 0.7170435      1 0.7996595

```

```
# ----- #
```

Examine baseline imbalance and conduct matching

We start by fitting a linear model to our data. Because `method = NULL`, no propensity score matching is done (i.e., baseline model)

```
# Because method = NULL, no propensity score matching
# is done (i.e., baseline model)
baseline <- matchit(treat ~ enrollment + st_ratio + susp + farms +
                    sped + minority + ell + disable + read,
                    data = df, method = NULL)
```


Let's examine standardized mean differences in covariates

```
bal.tab(baseline, s.d.denom = "treat", m.threshold = 0.1)
```

```
## Call
## matchit(formula = treat ~ enrollment + st_ratio + susp + farms +
##      sped + minority + ell + disable + read, data = df, method = NULL)
##
## Balance Measures
##      Type Diff.Un      M.Threshold.Un
## distance Distance 0.6102
## enrollment Contin. 0.0169      Balanced, <0.1
## st_ratio    Contin. 0.3444 Not Balanced, >0.1
## susp        Contin. 0.1335 Not Balanced, >0.1
## farms       Contin. 0.2918 Not Balanced, >0.1
## sped        Contin. 0.1065 Not Balanced, >0.1
## minority    Contin. 0.3799 Not Balanced, >0.1
## ell         Contin. 0.3122 Not Balanced, >0.1
## disable     Contin. 0.0416      Balanced, <0.1
## read        Contin. -0.1522 Not Balanced, >0.1
##
## Balance tally for mean differences
##      count
## Balanced, <0.1      2
## Not Balanced, >0.1  7
##
## Variable with the greatest mean difference
##      Variable Diff.Un      M.Threshold.Un
## minority 0.3799 Not Balanced, >0.1
##
## Sample sizes
##      Control Treated
## All      487      513
```

We can also plot our values.

```
# Plot SMDs
love.plot(baseline, s.d.denom = "pooled", abs = T,
          thresholds = 0.1, var.order = "unadjusted", drop.distance = T)
```



Now we will estimate propensity scores and match accordingly.

```
# Conduct propensity score matching
m <- matchit(treat ~ enrollment + st_ratio + susp + farms +
             sped + minority + ell + disable + read,
             data = df, method = "nearest", replace = T,
             ratio = 2, caliper = 0.1, std.caliper = T)
```

Examine our output: how do things look after matching?

```
summary(m)
```

```
##
## Call:
## matchit(formula = treat ~ enrollment + st_ratio + susp + farms +
##         sped + minority + ell + disable + read, data = df, method = "nearest",
##         replace = T, caliper = 0.1, std.caliper = T, ratio = 2)
##
```

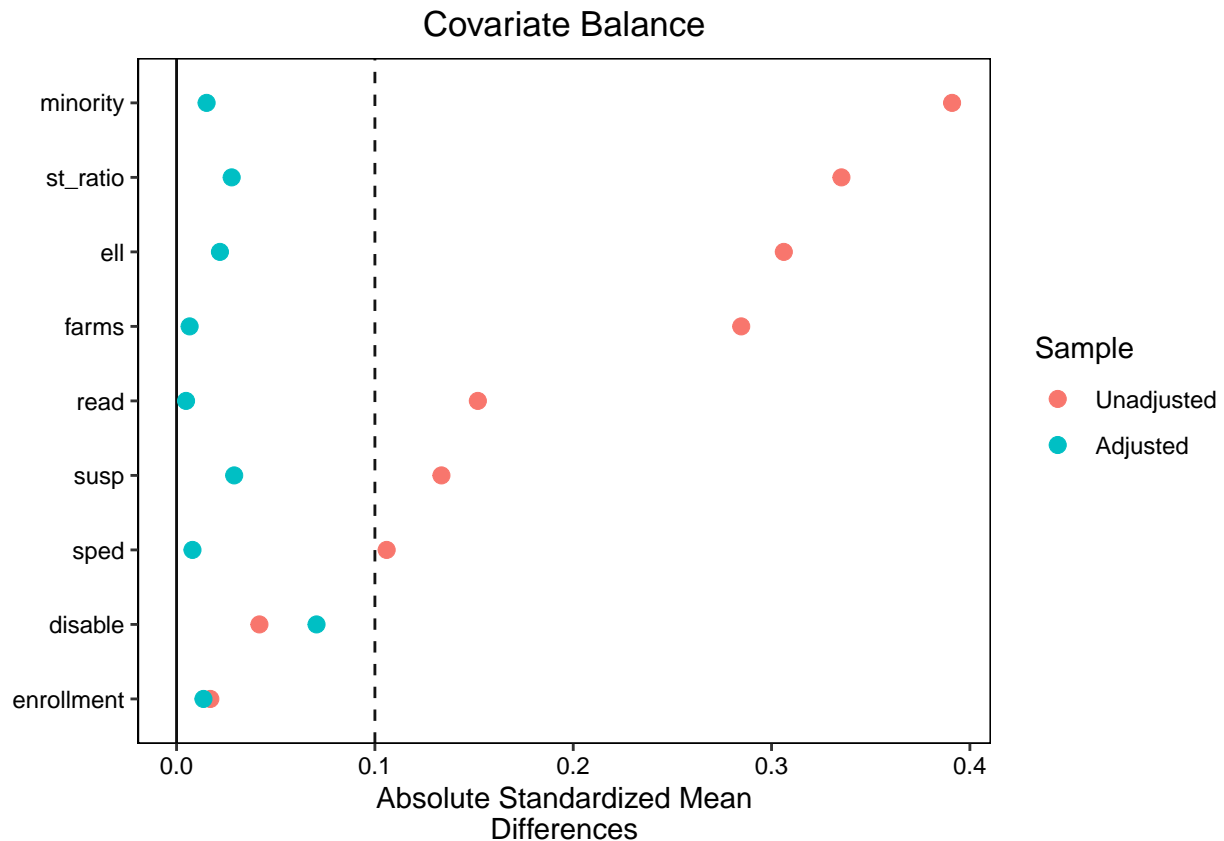
```

## Summary of Balance for All Data:
##           Means Treated Means Control Std. Mean Diff. Var. Ratio eCDF Mean
## distance           0.5526           0.4713           0.6102           0.9129           0.1650
## enrollment         558.2300          553.7639           0.0169           1.0231           0.0116
## st_ratio           0.2071           0.1914           0.3444           0.8998           0.0953
## susp              0.1030           0.0952           0.1335           1.0004           0.0391
## farms             0.2082           0.1919           0.2918           0.9079           0.0777
## sped             0.1512           0.1481           0.1065           0.9806           0.0316
## minority          0.4435           0.3561           0.3799           1.1258           0.1100
## ell              0.2087           0.1904           0.3122           0.9262           0.0887
## disable          0.1506           0.1494           0.0416           1.0201           0.0117
## read            0.7984           0.8062          -0.1522           0.9928           0.0339
##           eCDF Max
## distance          0.2767
## enrollment        0.0360
## st_ratio          0.1515
## susp             0.0760
## farms            0.1323
## sped            0.0829
## minority         0.2041
## ell             0.1540
## disable          0.0369
## read            0.0546
##
##
## Summary of Balance for Matched Data:
##           Means Treated Means Control Std. Mean Diff. Var. Ratio eCDF Mean
## distance           0.5505           0.5505          -0.0002           0.9985           0.0018
## enrollment         556.9470          553.3900           0.0135           1.0143           0.0111
## st_ratio           0.2068           0.2081          -0.0286           1.0133           0.0188
## susp              0.1029           0.1012           0.0291           0.9588           0.0220
## farms             0.2079           0.2083          -0.0067           0.9414           0.0126
## sped             0.1510           0.1507           0.0081           0.9151           0.0179
## minority          0.4414           0.4381           0.0147           1.1498           0.0277
## ell              0.2083           0.2096          -0.0223           0.9289           0.0117
## disable          0.1507           0.1486           0.0702           0.9673           0.0194
## read            0.7990           0.7992          -0.0048           0.9752           0.0205
##           eCDF Max Std. Pair Dist.
## distance          0.0118           0.0080
## enrollment        0.0422           1.1444
## st_ratio          0.0639           0.9364
## susp             0.0658           1.1200
## farms            0.0521           1.0198
## sped            0.0501           1.1161
## minority         0.0668           0.7967
## ell             0.0354           0.9958
## disable          0.0570           1.1115
## read            0.0521           1.1190
##
## Percent Balance Improvement:
##           Std. Mean Diff. Var. Ratio eCDF Mean eCDF Max
## distance           100.0           98.4           98.9           95.7
## enrollment          20.4           38.0           3.6          -17.5
## st_ratio            91.7           87.5           80.2           57.9

```

```
## susp          78.2   -11201.2    43.8    13.4
## farms         97.7     37.5     83.8    60.7
## sped          92.4   -353.1     43.2    39.6
## minority      96.1    -17.8     74.8    67.3
## ell           92.9     3.8     86.8    77.0
## disable       -68.9    -67.1    -65.7   -54.4
## read          96.9   -244.8     39.5     4.6
##
## Sample Sizes:
##           Control Treated
## All         487.     513
## Matched (ESS) 239.89   509
## Matched      363.     509
## Unmatched    124.       4
## Discarded     0.       0
```

```
love.plot(m, s.d.denom = "pooled", abs = T, thresholds = 0.1,
          var.order = "unadjusted", drop.distance = T)
```



The standardized mean differences decreased for every covariate except for 'disable' (although SMD < 0.1).

Estimate outcomes

Great - it looks like we have achieved balance on our covariates. Now let's estimate the treatment effect on math outcomes. We begin by retaining matched units only and including weights (as we matched with replacement).

```
match_data <- match.data(m)

# Then use weights, as matching with replacement
mwr_data <- svydesign(ids=~1, weights =~ weights, data = match_data)
```

Finally, we estimate the ATT (unadjusted such that treatment is the only predictor).

```
outcome_unadj <- svyglm(math ~ treat, mwr_data, family=gaussian())
```

Before we examine our estimates, the potential outcomes ATT = 0.0188. This can be interpreted as the true effect of treatment (as we generated the data and know both potential outcomes)

```
pop_ATT <- by(math_1, treat, mean)[[2]] - by(math_0, treat, mean)[[2]]
pop_ATT
```

```
## [1] 0.01888618
```

Now we can examine bias in our propensity score matched estimate of the treatment effect.

```
# Matching estimates
round(summary(outcome_unadj)$coef, digits = 4)
```

```
##           Estimate Std. Error  t value Pr(>|t|)
## (Intercept)   0.7954     0.0019 420.1098      0
## treat         0.0177     0.0023   7.6145      0
```

```
bias_outcome_unadj <- (coef(outcome_unadj)["treat"] - pop_ATT) / pop_ATT
bias_outcome_unadj
```

```
##           treat
## -0.06115827
```

Not bad, about 6.1% bias. Out of curiosity, what would our estimated treatment effect be had we not used propensity score matching, but fit a regression model in which we adjusted for covariates including treatment?

```
# Compare to estimates from basic regression adjustment
reg_adj <- lm(math ~ treat + enrollment + st_ratio + susp + farms +
              sped + minority + ell + disable + read, data = df)
round(summary(reg_adj)$coef, digits = 4)
```



```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.8030     0.0173 46.3652  0.0000
## treat        0.0167     0.0020  8.4388  0.0000
## enrollment   0.0000     0.0000  0.5974  0.5504
## st_ratio     -0.0153     0.0383 -0.4007  0.6887
## susp         0.0144     0.0165  0.8710  0.3840
## farms        0.0873     0.0581  1.5016  0.1335
## sped         0.0658     0.0319  2.0608  0.0396
## minority     0.0174     0.0043  4.0874  0.0000
## ell          -0.0240     0.0487 -0.4933  0.6219
## disable      0.0630     0.0328  1.9182  0.0554
## read         -0.0580     0.0188 -3.0860  0.0021
```

```
bias_reg_adj <- (coef(reg_adj)["treat"] - pop_ATT) / pop_ATT
bias_reg_adj
```

```
##      treat
## -0.1146654
```

```
# ----- #
```

While our PSM estimate is 6.1% biased, the regression adjustment estimate is 11.5% biased.